

Assignment 4

Adrian Bracher (Matr. Nr. 01637180)

21.04.2021

Contents

1	Transforming Data with dplyr	2
1.1	Selecting columns	2
1.2	Arranging observations	2
1.3	Filtering for conditions	2
1.4	Filtering and arranging	3
1.5	Calculating the number of government employees	3
1.6	Calculating the percentage of women in a county	4
1.7	Select, mutate, filter and arrange	4
2	Aggregating Data	5
2.1	Counting by region	5
2.2	Counting citizens by state	5
2.3	Mutating and counting	6
2.4	Summarizing	6
2.5	Summarizing by state	6
2.6	Summarizing by state and region	7
2.7	Selecting a county from each region	7
2.8	Finding the highest-income state in each region	8
2.9	Using summarize, top_n, and count together	8
3	Selecting and Transforming Data	8
3.1	Selecting columns	8
3.2	Select helpers	10
3.3	Renaming a column after count	10
3.4	Renaming a column as part of a select	10
3.5	Using transmute	11
3.6	Choosing among the four verbs	11
4	Case Study: The babynames Dataset	13
4.1	Filtering and arranging for one year	13
4.2	Using top_n with babynames	14
4.3	Visualizing names with ggplot2	14
4.4	Finding the year each name is most common	15
4.5	Adding the total and maximum for each name	16
4.6	Visualizing the normalized change in popularity	16
4.7	Using ratios to describe the frequency of a name	17
4.8	Biggest jumps in a name	18

1 Transforming Data with dplyr

1.1 Selecting columns

We select the columns state, county, population and poverty with a pipe syntax like below.

```
# Select the columns
counties %>%
  select(state, county, population, poverty)
```

```
## # A tibble: 3,138 x 4
##   state   county   population poverty
##   <chr>   <chr>         <dbl>   <dbl>
## 1 Alabama Autauga      55221    12.9
## 2 Alabama Baldwin     195121   13.4
## 3 Alabama Barbour     26932    26.7
## 4 Alabama Bibb        22604    16.8
## 5 Alabama Blount      57710    16.7
## 6 Alabama Bullock     10678    24.6
## 7 Alabama Butler      20354    25.4
## 8 Alabama Calhoun     116648    20.5
## 9 Alabama Chambers    34079    21.6
## 10 Alabama Cherokee   26008    19.2
## # ... with 3,128 more rows
```

1.2 Arranging observations

In this exercise we sort the selected columns in descending order on the column public_work.

```
counties_selected <- counties %>%
  select(state, county, population, private_work, public_work, self_employed)

# Add a verb to sort in descending order of public_work
counties_selected %>%
  arrange(desc(public_work))

## # A tibble: 3,138 x 6
##   state   county   population private_work public_work self_employed
##   <chr>   <chr>         <dbl>         <dbl>         <dbl>         <dbl>
## 1 Hawaii   Kalawao         85           25          64.1          10.9
## 2 Alaska   Yukon-Koyukuk Ce~ 5644         33.3          61.7           5.1
## 3 Wisconsin Menominee     4451         36.8          59.1           3.7
## 4 North Da~ Sioux      4380         32.9          56.8          10.2
## 5 South Da~ Todd      9942         34.4           55           9.8
## 6 Alaska   Lake and Peninsu~ 1474         42.2          51.6           6.1
## 7 Californ~ Lassen    32645         42.6          50.5           6.8
## 8 South Da~ Buffalo     2038         48.4          49.5           1.8
## 9 South Da~ Dewey     5579         34.9          49.2          14.7
## 10 Texas    Kenedy         565         51.9          48.1           0
## # ... with 3,128 more rows
```

1.3 Filtering for conditions

Filter() can be used like below to filter the data using one (or more) criteria.

```
counties_selected <- counties %>%
  select(state, county, population)
```

```
# Filter for counties in the state of California that have a population above 1000000
counties_selected %>%
  filter(state == "California", population > 1000000)
```

```
## # A tibble: 9 x 3
##   state      county      population
##   <chr>      <chr>      <dbl>
## 1 California Alameda      1584983
## 2 California Contra Costa  1096068
## 3 California Los Angeles  10038388
## 4 California Orange      3116069
## 5 California Riverside   2298032
## 6 California Sacramento   1465832
## 7 California San Bernardino 2094769
## 8 California San Diego     3223096
## 9 California Santa Clara   1868149
```

1.4 Filtering and arranging

The functions `arrange()` and `filter()` can also be used consecutively.

```
counties_selected <- counties %>%
  select(state, county, population, private_work, public_work, self_employed)

# Filter for Texas and more than 10000 people; sort in descending order of private_work
counties_selected %>%
  arrange(desc(private_work)) %>%
  filter(state == "Texas", population > 10000)
```

```
## # A tibble: 169 x 6
##   state county      population private_work public_work self_employed
##   <chr> <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Texas Gregg      123178        84.7        9.8        5.4
## 2 Texas Collin     862215        84.1       10        5.8
## 3 Texas Dallas   2485003        83.9        9.5        6.4
## 4 Texas Harris   4356362        83.4       10.1        6.3
## 5 Texas Andrews    16775        83.1        9.6        6.8
## 6 Texas Tarrant  1914526        83.1       11.4        5.4
## 7 Texas Titus      32553        82.5       10        7.4
## 8 Texas Denton    731851        82.2       11.9        5.7
## 9 Texas Ector     149557        82        11.2        6.7
## 10 Texas Moore     22281        82        11.7        5.9
## # ... with 159 more rows
```

1.5 Calculating the number of government employees

In this exercise we learn that `mutate` can be used to introduce a new column to the data frame and also that it can be filled with a calculation using the data in other columns.

```
counties_selected <- counties %>%
  select(state, county, population, public_work)

# Sort in descending order of the public_workers column
counties_selected %>%
```

```
mutate(public_workers = public_work * population / 100) %>%
  arrange(desc(public_workers))
```

```
## # A tibble: 3,138 x 5
##   state      county      population public_work public_workers
##   <chr>      <chr>          <dbl>      <dbl>      <dbl>
## 1 California Los Angeles    10038388      11.5    1154415.
## 2 Illinois   Cook              5236393      11.5     602185.
## 3 California San Diego      3223096      14.8     477018.
## 4 Arizona    Maricopa          4018143      11.7     470123.
## 5 Texas      Harris            4356362      10.1     439993.
## 6 New York   Kings             2595259      14.4     373717.
## 7 California San Bernardino  2094769      16.7     349826.
## 8 California Riverside     2298032      14.9     342407.
## 9 California Sacramento    1465832      21.8     319551.
## 10 California Orange       3116069      10.2     317839.
## # ... with 3,128 more rows
```

1.6 Calculating the percentage of women in a county

```
# Select the columns state, county, population, men, and women
counties_selected <- counties %>%
  select(state, county, population, men, women)
```

```
# Calculate proportion_women as the fraction of the population made up of women
counties_selected %>%
  mutate(proportion_women = women/population)
```

```
## # A tibble: 3,138 x 6
##   state      county      population  men women proportion_women
##   <chr>      <chr>          <dbl> <dbl> <dbl>      <dbl>
## 1 Alabama Autauga        55221 26745 28476      0.516
## 2 Alabama Baldwin      195121 95314 99807      0.512
## 3 Alabama Barbour      26932 14497 12435      0.462
## 4 Alabama Bibb          22604 12073 10531      0.466
## 5 Alabama Blount       57710 28512 29198      0.506
## 6 Alabama Bullock      10678  5660  5018      0.470
## 7 Alabama Butler       20354  9502 10852      0.533
## 8 Alabama Calhoun     116648 56274 60374      0.518
## 9 Alabama Chambers    34079 16258 17821      0.523
## 10 Alabama Cherokee    26008 12975 13033      0.501
## # ... with 3,128 more rows
```

1.7 Select, mutate, filter and arrange

In this exercise we combine everything from this chapter.

```
counties %>%
  # Select the five columns
  select(state, county, population, men, women) %>%
  # Add the proportion_men variable
  mutate(proportion_men = men/population) %>%
  # Filter for population of at least 10,000
```

```
filter(population > 10000) %>%
# Arrange proportion of men in descending order
arrange(desc(proportion_men))
```

```
## # A tibble: 2,437 x 6
##   state      county      population    men women proportion_men
##   <chr>     <chr>          <dbl> <dbl> <dbl>          <dbl>
## 1 Virginia Sussex          11864   8130  3734          0.685
## 2 California Lassen          32645  21818 10827          0.668
## 3 Georgia   Chattahoochee    11914   7940  3974          0.666
## 4 Louisiana West Feliciana  15415  10228  5187          0.664
## 5 Florida   Union           15191   9830  5361          0.647
## 6 Texas     Jones           19978  12652  7326          0.633
## 7 Missouri  DeKalb          12782   8080  4702          0.632
## 8 Texas     Madison         13838   8648  5190          0.625
## 9 Virginia  Greenville       11760   7303  4457          0.621
## 10 Texas    Anderson        57915  35469 22446          0.612
## # ... with 2,427 more rows
```

2 Aggregating Data

2.1 Counting by region

Count() is used here to count the number of regions.

```
# Use count to find the number of counties in each region
counties_selected %>%
  count(region, sort=TRUE)
```

2.2 Counting citizens by state

A weight can be applied too by specifying the parameter wt of the function count.

```
counties_selected <- counties
```

```
# Find number of counties per state, weighted by citizens
counties_selected %>%
  count(state, wt=citizens, sort = TRUE)
```

```
## # A tibble: 50 x 2
##   state      n
##   <chr>    <dbl>
## 1 California 24280349
## 2 Texas      16864864
## 3 Florida    13933052
## 4 New York   13531404
## 5 Pennsylvania 9710416
## 6 Illinois    8979999
## 7 Ohio        8709050
## 8 Michigan    7380136
## 9 North Carolina 7107998
## 10 Georgia     6978660
## # ... with 40 more rows
```

2.3 Mutating and counting

Mutate and count can be combined for advanced counting with weights.

```
counties_selected %>%  
  # Add population_walk containing the total number of people who walk to work  
  mutate(population_walk = population*walk/100) %>%  
  count(state, wt=population_walk, sort=TRUE)
```

```
## # A tibble: 50 x 2  
##   state      n  
##   <chr>    <dbl>  
## 1 New York 1237938.  
## 2 California 1017964.  
## 3 Pennsylvania 505397.  
## 4 Texas    430783.  
## 5 Illinois  400346.  
## 6 Massachusetts 316765.  
## 7 Florida   284723.  
## 8 New Jersey 273047.  
## 9 Ohio      266911.  
## 10 Washington 239764.  
## # ... with 40 more rows
```

2.4 Summarizing

Aggregations can be made with the function summarizing, for example calculating the min, max and mean like below.

```
# Summarize to find minimum population, maximum unemployment, and average income  
counties_selected %>%  
  summarize(min_population = min(population), max_unemployment=max(unemployment), average_income=mean(i
```

```
## # A tibble: 1 x 3  
##   min_population max_unemployment average_income  
##   <dbl>          <dbl>          <dbl>  
## 1      85          29.4          46832.
```

2.5 Summarizing by state

In the exercise below we calculate the population density for each state.

```
# Add a density column, then sort in descending order  
counties_selected %>%  
  group_by(state) %>%  
  summarize(total_area = sum(land_area),  
            total_population = sum(population)) %>%  
  mutate(density = total_population/total_area) %>%  
  arrange(desc(density))
```

```
## # A tibble: 50 x 4  
##   state      total_area total_population density  
##   <chr>          <dbl>          <dbl>    <dbl>  
## 1 New Jersey    7354.          8904413  1211.  
## 2 Rhode Island  1034.          1053661  1019.  
## 3 Massachusetts 7800.          6705586   860.  
## 4 Connecticut  4842.          3593222   742.
```

```
## 5 Maryland          9707.          5930538      611.
## 6 Delaware          1949.          926454       475.
## 7 New York          47126.         19673174      417.
## 8 Florida           53625.         19645772      366.
## 9 Pennsylvania      44743.         12779559      286.
## 10 Ohio             40861.         11575977      283.
## # ... with 40 more rows
```

2.6 Summarizing by state and region

In this exercise we learn that we can group by multiple columns at once.

```
# Calculate the average_pop and median_pop columns
```

```
counties_selected %>%
  group_by(region, state) %>%
  summarize(total_pop = sum(population)) %>%
  summarize(average_pop = mean(total_pop),
            median_pop = median(total_pop))
```

```
## `summarise()` has grouped output by 'region'. You can override using the `.groups` argument.
```

```
## # A tibble: 4 x 3
##   region      average_pop median_pop
##   <chr>          <dbl>      <dbl>
## 1 North Central  5627687.    5580644
## 2 Northeast     6221058.    3593222
## 3 South         7370486    4804098
## 4 West          5722755.    2798636
```

2.7 Selecting a county from each region

`top_n` is then used to select the `n` (the first parameter, in this case 1) rows with the most walking-to-work citizens.

```
# Group by region and find the greatest number of citizens who walk to work
```

```
counties_selected %>%
  group_by(region) %>%
  top_n(1, walk)
```

```
## # A tibble: 4 x 40
## # Groups:   region [4]
##   census_id state county region metro population men women hispanic white
##   <chr>      <chr> <chr> <chr> <chr>      <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 2013      Alaska Aleutia~ West Nonm~      3304  2198  1106    12    15
## 2 36061     New Y~ New York North~ Metro  1629507 769434 860073   25.8  47.1
## 3 38051     North~ McIntosh North~ Nonm~      2759  1341  1418    0.9   95.8
## 4 51678     Virgi~ Lexingt~ South Nonm~      7071  4372  2699    3.9   75.4
## # ... with 30 more variables: black <dbl>, native <dbl>, asian <dbl>,
## #   pacific <dbl>, citizens <dbl>, income <dbl>, income_err <dbl>,
## #   income_per_cap <dbl>, income_per_cap_err <dbl>, poverty <dbl>,
## #   child_poverty <dbl>, professional <dbl>, service <dbl>, office <dbl>,
## #   construction <dbl>, production <dbl>, drive <dbl>, carpool <dbl>,
## #   transit <dbl>, walk <dbl>, other_transp <dbl>, work_at_home <dbl>,
## #   mean_commute <dbl>, employed <dbl>, private_work <dbl>, public_work <dbl>,
## #   self_employed <dbl>, family_work <dbl>, unemployment <dbl>, land_area <dbl>
```

2.8 Finding the highest-income state in each region

In this exercise we select the highest-income state for each region.

```
counties_selected %>%
  group_by(region, state) %>%
  # Calculate average income
  summarize(average_income = mean(income)) %>%
  # Find the highest income state in each region
  top_n(1, average_income)
```

`summarise()` has grouped output by 'region'. You can override using the `.groups` argument.

```
## # A tibble: 4 x 3
## # Groups:   region [4]
##   region      state      average_income
##   <chr>      <chr>      <dbl>
## 1 North Central North Dakota    55575.
## 2 Northeast    New Jersey     73014.
## 3 South        Maryland      69200.
## 4 West         Alaska        65125.
```

2.9 Using summarize, top_n, and count together

In this exercise we combine everything we've learned in this chapter to count the number of states that have more people living in Metro areas and the ones that have more people living in Nonmetro areas.

```
# Count the states with more people in Metro or Nonmetro areas
counties_selected %>%
  group_by(state, metro) %>%
  summarize(total_pop = sum(population)) %>%
  top_n(1, total_pop) %>%
  ungroup() %>%
  count(metro)
```

`summarise()` has grouped output by 'state'. You can override using the `.groups` argument.

```
## # A tibble: 2 x 2
##   metro      n
##   <chr>   <int>
## 1 Metro      44
## 2 Nonmetro     6
```

3 Selecting and Transforming Data

3.1 Selecting columns

In the code below different columns are selected and then sorted by the column service.

```
# Glimpse the counties table
glimpse(counties)
```

```
## Rows: 3,138
## Columns: 40
## $ census_id      <chr> "1001", "1003", "1005", "1007", "1009", "1011", "10~
## $ state          <chr> "Alabama", "Alabama", "Alabama", "Alabama", "Alabam~
## $ county         <chr> "Autauga", "Baldwin", "Barbour", "Bibb", "Blount", ~
## $ region         <chr> "South", "South", "South", "South", "South", "South~
```



```
## $ metro <chr> "Metro", "Metro", "Nonmetro", "Metro", "Metro", "No~
## $ population <dbl> 55221, 195121, 26932, 22604, 57710, 10678, 20354, 1~
## $ men <dbl> 26745, 95314, 14497, 12073, 28512, 5660, 9502, 5627~
## $ women <dbl> 28476, 99807, 12435, 10531, 29198, 5018, 10852, 603~
## $ hispanic <dbl> 2.6, 4.5, 4.6, 2.2, 8.6, 4.4, 1.2, 3.5, 0.4, 1.5, 7~
## $ white <dbl> 75.8, 83.1, 46.2, 74.5, 87.9, 22.2, 53.3, 73.0, 57.~
## $ black <dbl> 18.5, 9.5, 46.7, 21.4, 1.5, 70.7, 43.8, 20.3, 40.3,~
## $ native <dbl> 0.4, 0.6, 0.2, 0.4, 0.3, 1.2, 0.1, 0.2, 0.2, 0.6, 0~
## $ asian <dbl> 1.0, 0.7, 0.4, 0.1, 0.1, 0.2, 0.4, 0.9, 0.8, 0.3, 0~
## $ pacific <dbl> 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0~
## $ citizens <dbl> 40725, 147695, 20714, 17495, 42345, 8057, 15581, 88~
## $ income <dbl> 51281, 50254, 32964, 38678, 45813, 31938, 32229, 41~
## $ income_err <dbl> 2391, 1263, 2973, 3995, 3141, 5884, 1793, 925, 2949~
## $ income_per_cap <dbl> 24974, 27317, 16824, 18431, 20532, 17580, 18390, 21~
## $ income_per_cap_err <dbl> 1080, 711, 798, 1618, 708, 2055, 714, 489, 1366, 15~
## $ poverty <dbl> 12.9, 13.4, 26.7, 16.8, 16.7, 24.6, 25.4, 20.5, 21.~
## $ child_poverty <dbl> 18.6, 19.2, 45.3, 27.9, 27.2, 38.4, 39.2, 31.6, 37.~
## $ professional <dbl> 33.2, 33.1, 26.8, 21.5, 28.5, 18.8, 27.5, 27.3, 23.~
## $ service <dbl> 17.0, 17.7, 16.1, 17.9, 14.1, 15.0, 16.6, 17.7, 14.~
## $ office <dbl> 24.2, 27.1, 23.1, 17.8, 23.9, 19.7, 21.9, 24.2, 26.~
## $ construction <dbl> 8.6, 10.8, 10.8, 19.0, 13.5, 20.1, 10.3, 10.5, 11.5~
## $ production <dbl> 17.1, 11.2, 23.1, 23.7, 19.9, 26.4, 23.7, 20.4, 24.~
## $ drive <dbl> 87.5, 84.7, 83.8, 83.2, 84.9, 74.9, 84.5, 85.3, 85.~
## $ carpool <dbl> 8.8, 8.8, 10.9, 13.5, 11.2, 14.9, 12.4, 9.4, 11.9, ~
## $ transit <dbl> 0.1, 0.1, 0.4, 0.5, 0.4, 0.7, 0.0, 0.2, 0.2, 0.2, 0~
## $ walk <dbl> 0.5, 1.0, 1.8, 0.6, 0.9, 5.0, 0.8, 1.2, 0.3, 0.6, 1~
## $ other_transp <dbl> 1.3, 1.4, 1.5, 1.5, 0.4, 1.7, 0.6, 1.2, 0.4, 0.7, 1~
## $ work_at_home <dbl> 1.8, 3.9, 1.6, 0.7, 2.3, 2.8, 1.7, 2.7, 2.1, 2.5, 1~
## $ mean_commute <dbl> 26.5, 26.4, 24.1, 28.8, 34.9, 27.5, 24.6, 24.1, 25.~
## $ employed <dbl> 23986, 85953, 8597, 8294, 22189, 3865, 7813, 47401,~
## $ private_work <dbl> 73.6, 81.5, 71.8, 76.8, 82.0, 79.5, 77.4, 74.1, 85.~
## $ public_work <dbl> 20.9, 12.3, 20.8, 16.1, 13.5, 15.1, 16.2, 20.8, 12.~
## $ self_employed <dbl> 5.5, 5.8, 7.3, 6.7, 4.2, 5.4, 6.2, 5.0, 2.8, 7.9, 4~
## $ family_work <dbl> 0.0, 0.4, 0.1, 0.4, 0.4, 0.0, 0.2, 0.1, 0.0, 0.5, 0~
## $ unemployment <dbl> 7.6, 7.5, 17.6, 8.3, 7.7, 18.0, 10.9, 12.3, 8.9, 7.~
## $ land_area <dbl> 594.44, 1589.78, 884.88, 622.58, 644.78, 622.81, 77~
```

```
counties %>%
```

```
# Select state, county, population, and industry-related columns
```

```
select(state, county, population, professional, service, office, construction, production) %>%
```

```
# Arrange service in descending order
```

```
arrange(desc(service))
```

```
## # A tibble: 3,138 x 8
```

	state	county	population	professional	service	office	construction	production
	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	Missis~	Tunica	10477	23.9	36.6	21.5	3.5	14.5
## 2	Texas	Kinney	3577	30	36.5	11.6	20.5	1.3
## 3	Texas	Kenedy	565	24.9	34.1	20.5	20.5	0
## 4	New Yo~	Bronx	1428357	24.3	33.3	24.2	7.1	11
## 5	Texas	Brooks	7221	19.6	32.4	25.3	11.1	11.5
## 6	Colora~	Fremo~	46809	26.6	32.2	22.8	10.7	7.6
## 7	Texas	Culbe~	2296	20.1	32.2	24.2	15.7	7.8
## 8	Califo~	Del N~	27788	33.9	31.5	18.8	8.9	6.8
## 9	Minnes~	Mahno~	5496	26.8	31.5	18.7	13.1	9.9

```
## 10 Virgin~ Lanca~      11129      30.3   31.2   22.8      8.1      7.6
## # ... with 3,128 more rows
```

3.2 Select helpers

We learn that columns are not only selectable by their exact name, but also with helpers such as `ends_with()` and `starts_with()`.

```
counties %>%
  # Select the state, county, population, and those ending with "work"
  select(state, county, population, ends_with("work")) %>%
  # Filter for counties that have at least 50% of people engaged in public work
  filter(public_work >= 50)
```

```
## # A tibble: 7 x 6
##   state      county      population private_work public_work family_work
##   <chr>    <chr>          <dbl>         <dbl>      <dbl>      <dbl>
## 1 Alaska  Lake and Peninsula~    1474         42.2       51.6        0.2
## 2 Alaska  Yukon-Koyukuk Cens~    5644         33.3       61.7         0
## 3 California Lassen            32645         42.6       50.5        0.1
## 4 Hawaii   Kalawao              85          25        64.1         0
## 5 North Dak~ Sioux            4380         32.9       56.8        0.1
## 6 South Dak~ Todd            9942         34.4        55         0.8
## 7 Wisconsin Menominee          4451         36.8       59.1        0.4
```

3.3 Renaming a column after count

The function `rename` can be used to rename a column, for example the default column name `n` of `count` to `num_counties`.

```
# Rename the n column to num_counties
counties %>%
  count(state) %>%
  rename(num_counties = n)
```

```
## # A tibble: 50 x 2
##   state      num_counties
##   <chr>          <int>
## 1 Alabama           67
## 2 Alaska            28
## 3 Arizona           15
## 4 Arkansas          75
## 5 California        58
## 6 Colorado          64
## 7 Connecticut         8
## 8 Delaware           3
## 9 Florida           67
## 10 Georgia          159
## # ... with 40 more rows
```

3.4 Renaming a column as part of a select

Renaming can also be done with `select`, below we rename the `poverty` to `poverty_rate`.

```
# Select state, county, and poverty as poverty_rate
counties %>%
  select(state, county, poverty_rate = poverty)
```

```
## # A tibble: 3,138 x 3
##   state county poverty_rate
##   <chr>  <chr>      <dbl>
## 1 Alabama Autauga      12.9
## 2 Alabama Baldwin      13.4
## 3 Alabama Barbour      26.7
## 4 Alabama Bibb         16.8
## 5 Alabama Blount       16.7
## 6 Alabama Bullock      24.6
## 7 Alabama Butler       25.4
## 8 Alabama Calhoun      20.5
## 9 Alabama Chambers     21.6
## 10 Alabama Cherokee    19.2
## # ... with 3,128 more rows
```

3.5 Using transmute

The verb `transmute` can be viewed as a combination of the `select` and `mutate` functionality. For an example view the code below.

```
counties %>%
  # Keep the state, county, and populations columns, and add a density column
  transmute(state, county, population, density = population/land_area) %>%
  # Filter for counties with a population greater than one million
  filter(population > 1000000) %>%
  # Sort density in ascending order
  arrange(density)
```

```
## # A tibble: 41 x 4
##   state county population density
##   <chr>  <chr>      <dbl>  <dbl>
## 1 California San Bernardino 2094769 104.
## 2 Nevada Clark 2035572 258.
## 3 California Riverside 2298032 319.
## 4 Arizona Maricopa 4018143 437.
## 5 Florida Palm Beach 1378806 700.
## 6 California San Diego 3223096 766.
## 7 Washington King 2045756 967.
## 8 Texas Travis 1121645 1133.
## 9 Florida Hillsborough 1302884 1277.
## 10 Florida Orange 1229039 1360.
## # ... with 31 more rows
```

3.6 Choosing among the four verbs

There are different use cases for `rename`, `select`, `mutate` and `transmute`: - `Rename` and `mutate` leave the columns that you don't mention, `select` and `transmute` don't. - `Rename` and `select` don't allow calculation, `mutate` and `transmute` do.

```
# Change the name of the unemployment column
counties %>%
  rename(unemployment_rate = unemployment)
```

```
## # A tibble: 3,138 x 40
```

```
## census_id state county region metro population men women hispanic white
## <chr> <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1001 Alabama Autauga South Metro 55221 26745 28476 2.6 75.8
## 2 1003 Alabama Baldwin South Metro 195121 95314 99807 4.5 83.1
## 3 1005 Alabama Barbour South Nonme~ 26932 14497 12435 4.6 46.2
## 4 1007 Alabama Bibb South Metro 22604 12073 10531 2.2 74.5
## 5 1009 Alabama Blount South Metro 57710 28512 29198 8.6 87.9
## 6 1011 Alabama Bullock South Nonme~ 10678 5660 5018 4.4 22.2
## 7 1013 Alabama Butler South Nonme~ 20354 9502 10852 1.2 53.3
## 8 1015 Alabama Calhoun South Metro 116648 56274 60374 3.5 73
## 9 1017 Alabama Chambe~ South Nonme~ 34079 16258 17821 0.4 57.3
## 10 1019 Alabama Chero~ South Nonme~ 26008 12975 13033 1.5 91.7
## # ... with 3,128 more rows, and 30 more variables: black <dbl>, native <dbl>,
## # asian <dbl>, pacific <dbl>, citizens <dbl>, income <dbl>, income_err <dbl>,
## # income_per_cap <dbl>, income_per_cap_err <dbl>, poverty <dbl>,
## # child_poverty <dbl>, professional <dbl>, service <dbl>, office <dbl>,
## # construction <dbl>, production <dbl>, drive <dbl>, carpool <dbl>,
## # transit <dbl>, walk <dbl>, other_transp <dbl>, work_at_home <dbl>,
## # mean_commute <dbl>, employed <dbl>, private_work <dbl>, public_work <dbl>,
## # self_employed <dbl>, family_work <dbl>, unemployment_rate <dbl>,
## # land_area <dbl>
```

```
# Keep the state and county columns, and the columns containing poverty
counties %>%
  select(state, county, contains("poverty"))
```

```
## # A tibble: 3,138 x 4
## state county poverty child_poverty
## <chr> <chr> <dbl> <dbl>
## 1 Alabama Autauga 12.9 18.6
## 2 Alabama Baldwin 13.4 19.2
## 3 Alabama Barbour 26.7 45.3
## 4 Alabama Bibb 16.8 27.9
## 5 Alabama Blount 16.7 27.2
## 6 Alabama Bullock 24.6 38.4
## 7 Alabama Butler 25.4 39.2
## 8 Alabama Calhoun 20.5 31.6
## 9 Alabama Chambers 21.6 37.2
## 10 Alabama Cherokee 19.2 30.1
## # ... with 3,128 more rows
```

```
# Calculate the fraction_women column without dropping the other columns
counties %>%
  mutate(fraction_women = women / population)
```

```
## # A tibble: 3,138 x 41
## census_id state county region metro population men women hispanic white
## <chr> <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1001 Alabama Autauga South Metro 55221 26745 28476 2.6 75.8
## 2 1003 Alabama Baldwin South Metro 195121 95314 99807 4.5 83.1
## 3 1005 Alabama Barbour South Nonme~ 26932 14497 12435 4.6 46.2
## 4 1007 Alabama Bibb South Metro 22604 12073 10531 2.2 74.5
## 5 1009 Alabama Blount South Metro 57710 28512 29198 8.6 87.9
## 6 1011 Alabama Bullock South Nonme~ 10678 5660 5018 4.4 22.2
## 7 1013 Alabama Butler South Nonme~ 20354 9502 10852 1.2 53.3
## 8 1015 Alabama Calhoun South Metro 116648 56274 60374 3.5 73
```

```
## 9 1017 Alabama Chamber~ South Nonme~ 34079 16258 17821 0.4 57.3
## 10 1019 Alabama Cheroke~ South Nonme~ 26008 12975 13033 1.5 91.7
## # ... with 3,128 more rows, and 31 more variables: black <dbl>, native <dbl>,
## # asian <dbl>, pacific <dbl>, citizens <dbl>, income <dbl>, income_err <dbl>,
## # income_per_cap <dbl>, income_per_cap_err <dbl>, poverty <dbl>,
## # child_poverty <dbl>, professional <dbl>, service <dbl>, office <dbl>,
## # construction <dbl>, production <dbl>, drive <dbl>, carpool <dbl>,
## # transit <dbl>, walk <dbl>, other_transp <dbl>, work_at_home <dbl>,
## # mean_commute <dbl>, employed <dbl>, private_work <dbl>, public_work <dbl>,
## # self_employed <dbl>, family_work <dbl>, unemployment <dbl>,
## # land_area <dbl>, fraction_women <dbl>

# Keep only the state, county, and employment_rate columns
counties %>%
  transmute(state, county, employment_rate = employed / population)

## # A tibble: 3,138 x 3
##   state county employment_rate
##   <chr> <chr>           <dbl>
## 1 Alabama Autauga 0.434
## 2 Alabama Baldwin 0.441
## 3 Alabama Barbour 0.319
## 4 Alabama Bibb 0.367
## 5 Alabama Blount 0.384
## 6 Alabama Bullock 0.362
## 7 Alabama Butler 0.384
## 8 Alabama Calhoun 0.406
## 9 Alabama Chambers 0.402
## 10 Alabama Cherokee 0.390
## # ... with 3,128 more rows
```

4 Case Study: The babynames Dataset

4.1 Filtering and arranging for one year

We use filter and arrange to select rows with year = 1990 and sort them by the number column. This results in the most common names for each year.

```
babynames %>%
  # Filter for the year 1990
  filter(year == 1990) %>%
  # Sort the number column in descending order
  arrange(desc(number))
```

```
## # A tibble: 21,223 x 3
##   year name      number
##   <dbl> <chr>      <int>
## 1 1990 Michael 65560
## 2 1990 Christopher 52520
## 3 1990 Jessica 46615
## 4 1990 Ashley 45797
## 5 1990 Matthew 44925
## 6 1990 Joshua 43382
## 7 1990 Brittany 36650
## 8 1990 Amanda 34504
```

```
## 9 1990 Daniel      33963
## 10 1990 David      33862
## # ... with 21,213 more rows
```

4.2 Using top_n with babynames

The code below results in the most popular name for each year.

```
# Find the most common name in each year
```

```
babynames %>%
  group_by(year) %>%
  top_n(1, number)
```

```
## # A tibble: 28 x 3
## # Groups:   year [28]
##   year name  number
##   <dbl> <chr> <int>
## 1 1880 John   9701
## 2 1885 Mary   9166
## 3 1890 Mary  12113
## 4 1895 Mary  13493
## 5 1900 Mary  16781
## 6 1905 Mary  16135
## 7 1910 Mary  22947
## 8 1915 Mary  58346
## 9 1920 Mary  71175
## 10 1925 Mary  70857
## # ... with 18 more rows
```

4.3 Visualizing names with ggplot2


The %in% keyword checks if a value is found in a vector. Then we plot number over year for the three names selected.

```
# Filter for the names Steven, Thomas, and Matthew
```

```
selected_names <- babynames %>%
  filter(name %in% c("Steven", "Thomas", "Matthew"))
```

```
# Plot the names using a different color for each name
```

```
ggplot(selected_names, aes(x = year, y = number, color = name)) +
  geom_line()
```



4.4 Finding the year each name is most common

We combine everything we've learned to find the most common baby name for each year.

Calculate the fraction of people born each year with the same name

```
babynames %>%
  group_by(year) %>%
  mutate(year_total = sum(number)) %>%
  ungroup() %>%
  mutate(fraction = number / year_total) %>%
```

Find the year each name is most common

```
group_by(name) %>%
  top_n(1, fraction)
```

A tibble: 48,040 x 5

Groups: name [48,040]

##	year	name	number	year_total	fraction
##	<dbl>	<chr>	<int>	<int>	<dbl>
## 1	1880	Abbott	5	201478	0.0000248
## 2	1880	Abe	50	201478	0.000248
## 3	1880	Abner	27	201478	0.000134
## 4	1880	Adelbert	28	201478	0.000139
## 5	1880	Adella	26	201478	0.000129
## 6	1880	Adolf	6	201478	0.0000298
## 7	1880	Adolph	93	201478	0.000462
## 8	1880	Agustus	5	201478	0.0000248
## 9	1880	Albert	1493	201478	0.00741

```
## 10 1880 Albertina      7      201478 0.0000347
## # ... with 48,030 more rows
```

4.5 Adding the total and maximum for each name

In this exercise we calculate the total occurrences of one name as `name_total` and also the max occurrence of one name `name_max`. Then we add a column `fraction_max` with `mutate`.

```
babynames %>%
  group_by(name) %>%
  mutate(name_total = sum(number),
         name_max = max(number)) %>%
  # Ungroup the table
  ungroup() %>%
  # Add the fraction_max column containing the number by the name maximum
  mutate(fraction_max = number/name_max)
```


```
## # A tibble: 332,595 x 6
##   year name      number name_total name_max fraction_max
##   <dbl> <chr>      <int>      <int>      <int>      <dbl>
## 1 1880 Aaron        102    114739    14635    0.00697
## 2 1880 Ab           5         77         31    0.161
## 3 1880 Abbie        71      4330      445    0.160
## 4 1880 Abbott        5       217        51    0.0980
## 5 1880 Abby         6     11272     1753    0.00342
## 6 1880 Abe         50      1832      271    0.185
## 7 1880 Abel         9     10565     3245    0.00277
## 8 1880 Abigail      12     72600    15762    0.000761
## 9 1880 Abner        27      1552      199    0.136
## 10 1880 Abraham     81     17882     2449    0.0331
## # ... with 332,585 more rows
```

4.6 Visualizing the normalized change in popularity

In this exercise we filter for three names and then plot the `fraction_max` column of the previous exercise with `ggplot`.

```
# Filter for the names Steven, Thomas, and Matthew
names_filtered <- names_normalized %>%
  filter(name %in% c("Steven", "Thomas", "Matthew"))

# Visualize these names over time
ggplot(names_filtered, aes(x=year, y=fraction_max, color=name)) + geom_line()
```

4.7 Using ratios to describe the frequency of a name

In this exercise we learn that the function `lag` can be used to find the previous value in vector/data frame.

```
babynames_fraction %>%
  # Arrange the data in order of name, then year
  arrange(name, year) %>%
  # Group the data by name
  group_by(name) %>%
  # Add a ratio column that contains the ratio between each year
  mutate(ratio = fraction / lag(fraction))
```

```
## # A tibble: 332,595 x 6
## # Groups:   name [48,040]
##   year name    number year_total fraction ratio
##   <dbl> <chr>    <int>    <int>    <dbl> <dbl>
## 1 2010 Aaban         9    3672066 0.00000245 NA
## 2 2015 Aaban        15    3648781 0.00000411 1.68
## 3 1995 Aadam         6    3652750 0.00000164 NA
## 4 2000 Aadam         6    3767293 0.00000159 0.970
## 5 2005 Aadam         6    3828460 0.00000157 0.984
## 6 2010 Aadam         7    3672066 0.00000191 1.22
## 7 2015 Aadam        22    3648781 0.00000603 3.16
## 8 2010 Aadan        11    3672066 0.00000300 NA
## 9 2015 Aadan        10    3648781 0.00000274 0.915
## 10 2000 Aadarsh      5    3767293 0.00000133 NA
## # ... with 332,585 more rows
```

4.8 Biggest jumps in a name

In this exercise we evaluate which names had the biggest jumps in popularity in consecutive years.

```
babynames_ratios_filtered %>%  
  # Extract the largest ratio from each name  
  top_n(1, ratio) %>%  
  # Sort the ratio column in descending order  
  arrange(desc(ratio)) %>%  
  # Filter for fractions greater than or equal to 0.001  
  filter(fraction >= 0.001)
```

```
## # A tibble: 291 x 6  
## # Groups:   name [291]  
##   year name      number year_total fraction ratio  
##   <dbl> <chr>    <int>      <int>    <dbl> <dbl>  
## 1 1960 Tammy     14365    4152075 0.00346 70.1  
## 2 2005 Nevaeh    4610    3828460 0.00120 45.8  
## 3 1940 Brenda    5460    2301630 0.00237 37.5  
## 4 1885 Grover     774     240822 0.00321 36.0  
## 5 1945 Cheryl    8170    2652029 0.00308 24.9  
## 6 1955 Lori      4980    4012691 0.00124 23.2  
## 7 2010 Khloe     5411    3672066 0.00147 23.2  
## 8 1950 Debra     6189    3502592 0.00177 22.6  
## 9 2010 Bentley   4001    3672066 0.00109 22.4  
## 10 1935 Marlene  4840    2088487 0.00232 16.8  
## # ... with 281 more rows
```